

Text Embeddings for Retrieval

Embeddings map text into dense vector spaces where semantic similarity corresponds to geometric closeness.

Key points:

- Vectors are typically 384 to 3072 dimensions.
- Cosine similarity is used to rank retrieved chunks.
- Embeddings are computed once at ingestion and stored in a vector index.
- Chunking strategy strongly affects retrieval quality.